

AI-Powered Anomaly Detection: Spotting Anomalies in Real-Time

Joe Khazen

Principal, WW GTM Specialist,
Data Streaming (AWS)

Austin Groeneveld

Assoc. WW SSA Data
Streaming (AWS)

Ali Alemi

Sr. WW SSA Data
Streaming (AWS)



Goals for the Day

- Get hands-on experience with multiple AWS services to identify nefarious behavior and anomalies in network flow logs in real-time
- Diagnose and recommend actions
- Use Amazon Bedrock for integrating Generative AI
- Send e-mail notification to interested parties
- Network
- Have fun

Services used today

- **Amazon MSK- Amazon Managed Service for Apache Kafka**
- **ADF-Amazon Data Firehose**
- **Amazon MSF-Amazon Managed Service for Apache Flink**
- **Amazon OpenSearch**
- **Amazon SageMaker**
- **Amazon Bedrock**
- **Amazon SNS-Simple Notification Service**

Agenda

Introduction

- Discuss the business case and AWS Services Overview

Ingest and Examine Real-time Data

- Produce and validate simulated data.

Processing Data in Real-time and Create Smart Alerts with Generative AI

- Validate model in Sagemaker, start a durable MSF Application with state that queries Sagemaker in real-time, validate the output of the MSF application in MSF Studio, all while interacting with MSK Serverless. Explore LLMs on Amazon Bedrock and the Amazon SNS topic for pushing alerts.

Configure Amazon Data Firehose and OpenSearch Ingestion

- Deliver data to S3 using Amazon Data Firehose and to Amazon OpenSearch for further exploration and analysis.

Guidelines for the session

- **Stay on Mute:** Keep your microphone muted when not speaking to avoid background noise.
- **Raise Your Hand for Help in Chime:** If you need assistance, please use the “Raise Hand” feature, or type “Help” in the chat. A facilitator will assist you as soon as possible.
- **Ask Questions in the Chime Chat:** For general questions, please use the chat box. This allows us to address questions without interrupting the flow of the lab.
- **Stay for the Full Lab:** Please commit to staying for the entire session. Hopping in and out disrupts the experience for others and may cause you to miss key steps.
- **Follow Along and Participate:** This is a hands-on lab, so we encourage you to actively follow along with the exercises and avoid multitasking.




Data Streaming and Anomaly Detection

- Why? What **problems** are we solving?
- Architecture
- Overview of **Solution Components**
- Workshop (**with Checkpoints**)

Streaming data technologies can unlock unprecedented value for organizations today




**Ingest from
various sources**



**Optimize and
reduce costs**




**Faster insights and
better analytics**



**Keep your systems
updated and
synchronized**


Value of Streaming Data for Anomaly Detection



Use Streaming and AI to reduce
time to detect & time to
respond

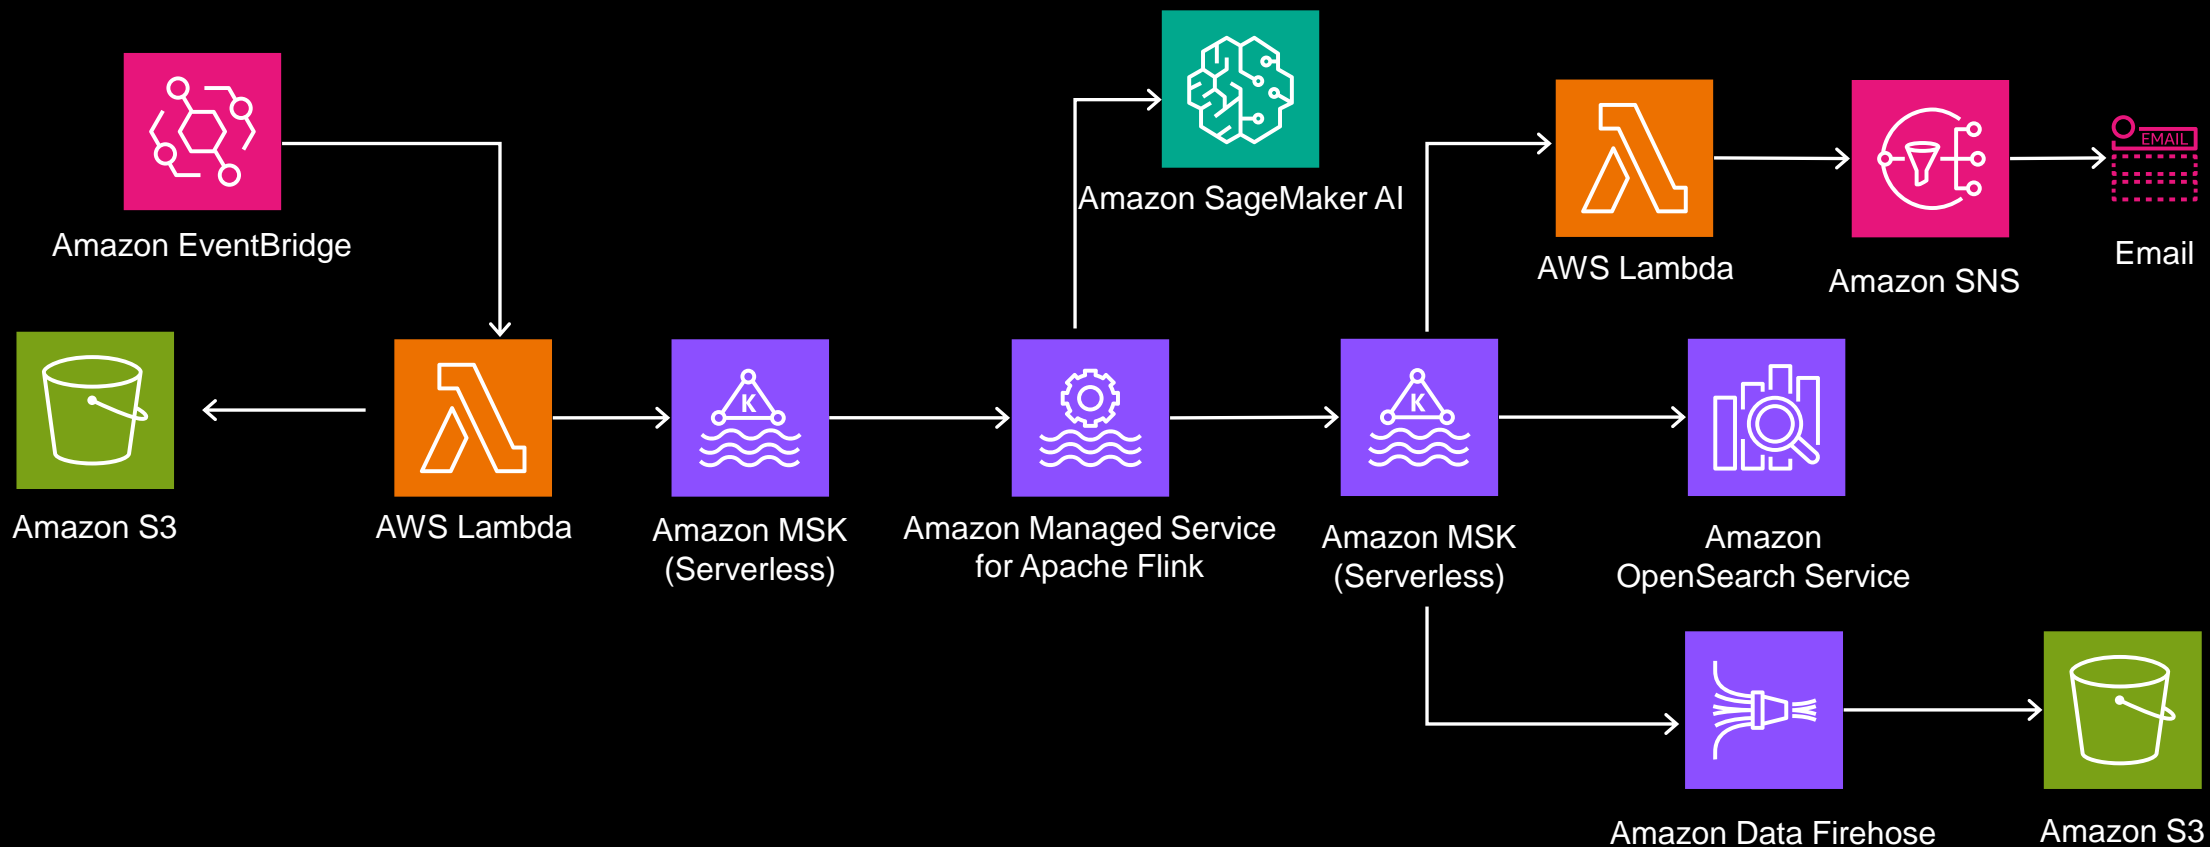


Continuous training the model
new traffic patterns

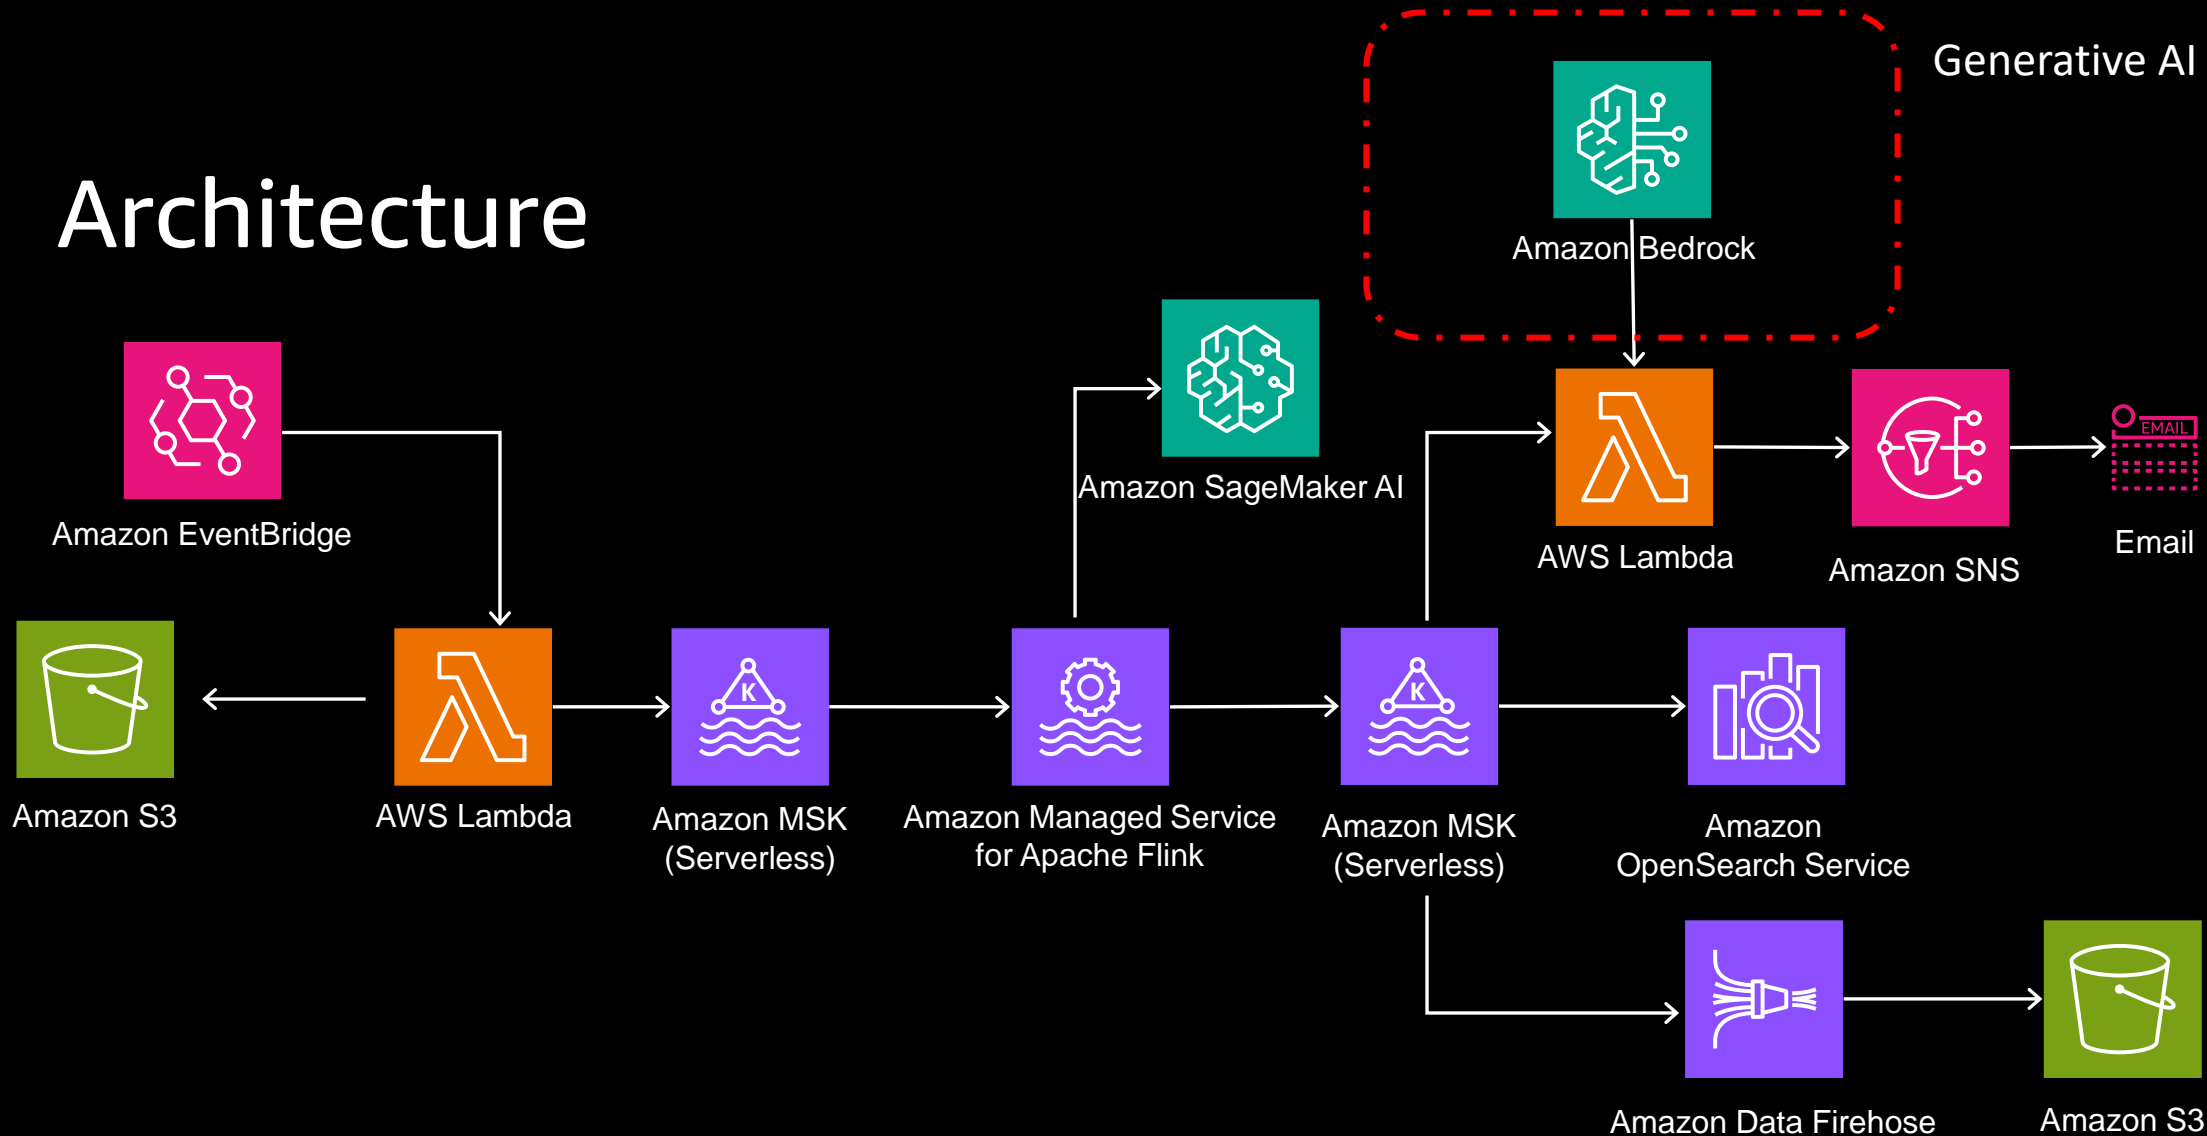


Simplifying Integrations
across various components

Architecture



Architecture




Challenge Statement

Businesses would like to leverage their existing network flow logs to identify nefarious behavior and attacks on their network

Your objectives:

- Identify nefarious behavior
- Diagnose and recommend actions
- Send e-mail notification to interested parties

... in real-time.




How?

- Streaming Analytics leveraging open source software
 - Apache Kafka (Amazon Managed Streaming for Apache Kafka)
 - Apache Flink (Amazon Managed Service for Apache Flink)
- Detect Anomalies using Amazon Sagemaker AI
- Identify and recommend next steps using Amazon Bedrock
- And more – Lambda, SNS, S3, OpenSearch, if you are interested

Informal Poll

Familiar with Apache Kafka?



Familiar with Apache Flink?



KAFKA is a registered trademark of The Apache Software Foundation and has been licensed for use by AWS. AWS has no affiliation with and is not endorsed by The Apache Software Foundation.

Apache Kafka 101

Apache Kafka is a “Distributed Streaming Platform”

Apache Kafka has 3 core capabilities:

- Allows applications to publish and subscribe to streams of records
- Store records in the same way they were received
- Allows applications to process records as they occurred

What is Apache Flink ?



A powerful open-source framework and engine for processing data streams



Diverse use-cases

- Process both batch and streaming data
- Transform and act on streaming data
- Extract value from streaming data
- Event-driven applications
- Streaming Analytics &ETL
- Batch analytics

Amazon SageMaker AI

"End-to-End Machine Learning Platform"

- **Data Preparation:** Label, Prepare, and Process your machine learning data
- **Model Training:** Train models using built-in algorithms or custom containers.
- **Model Deployment:** Deploy models to production with auto-scaling and real-time inference.
- **MLOps:** Manage the entire machine learning lifecycle, including monitoring and governance, with integrated tools.

Amazon Bedrock

"Foundation Model Service for Generative AI"

Pre-built Foundation Models: Access pre-trained foundation models from AWS, AI21 Labs, Anthropic, Cohere, Meta, Mistral AI, Stability AI

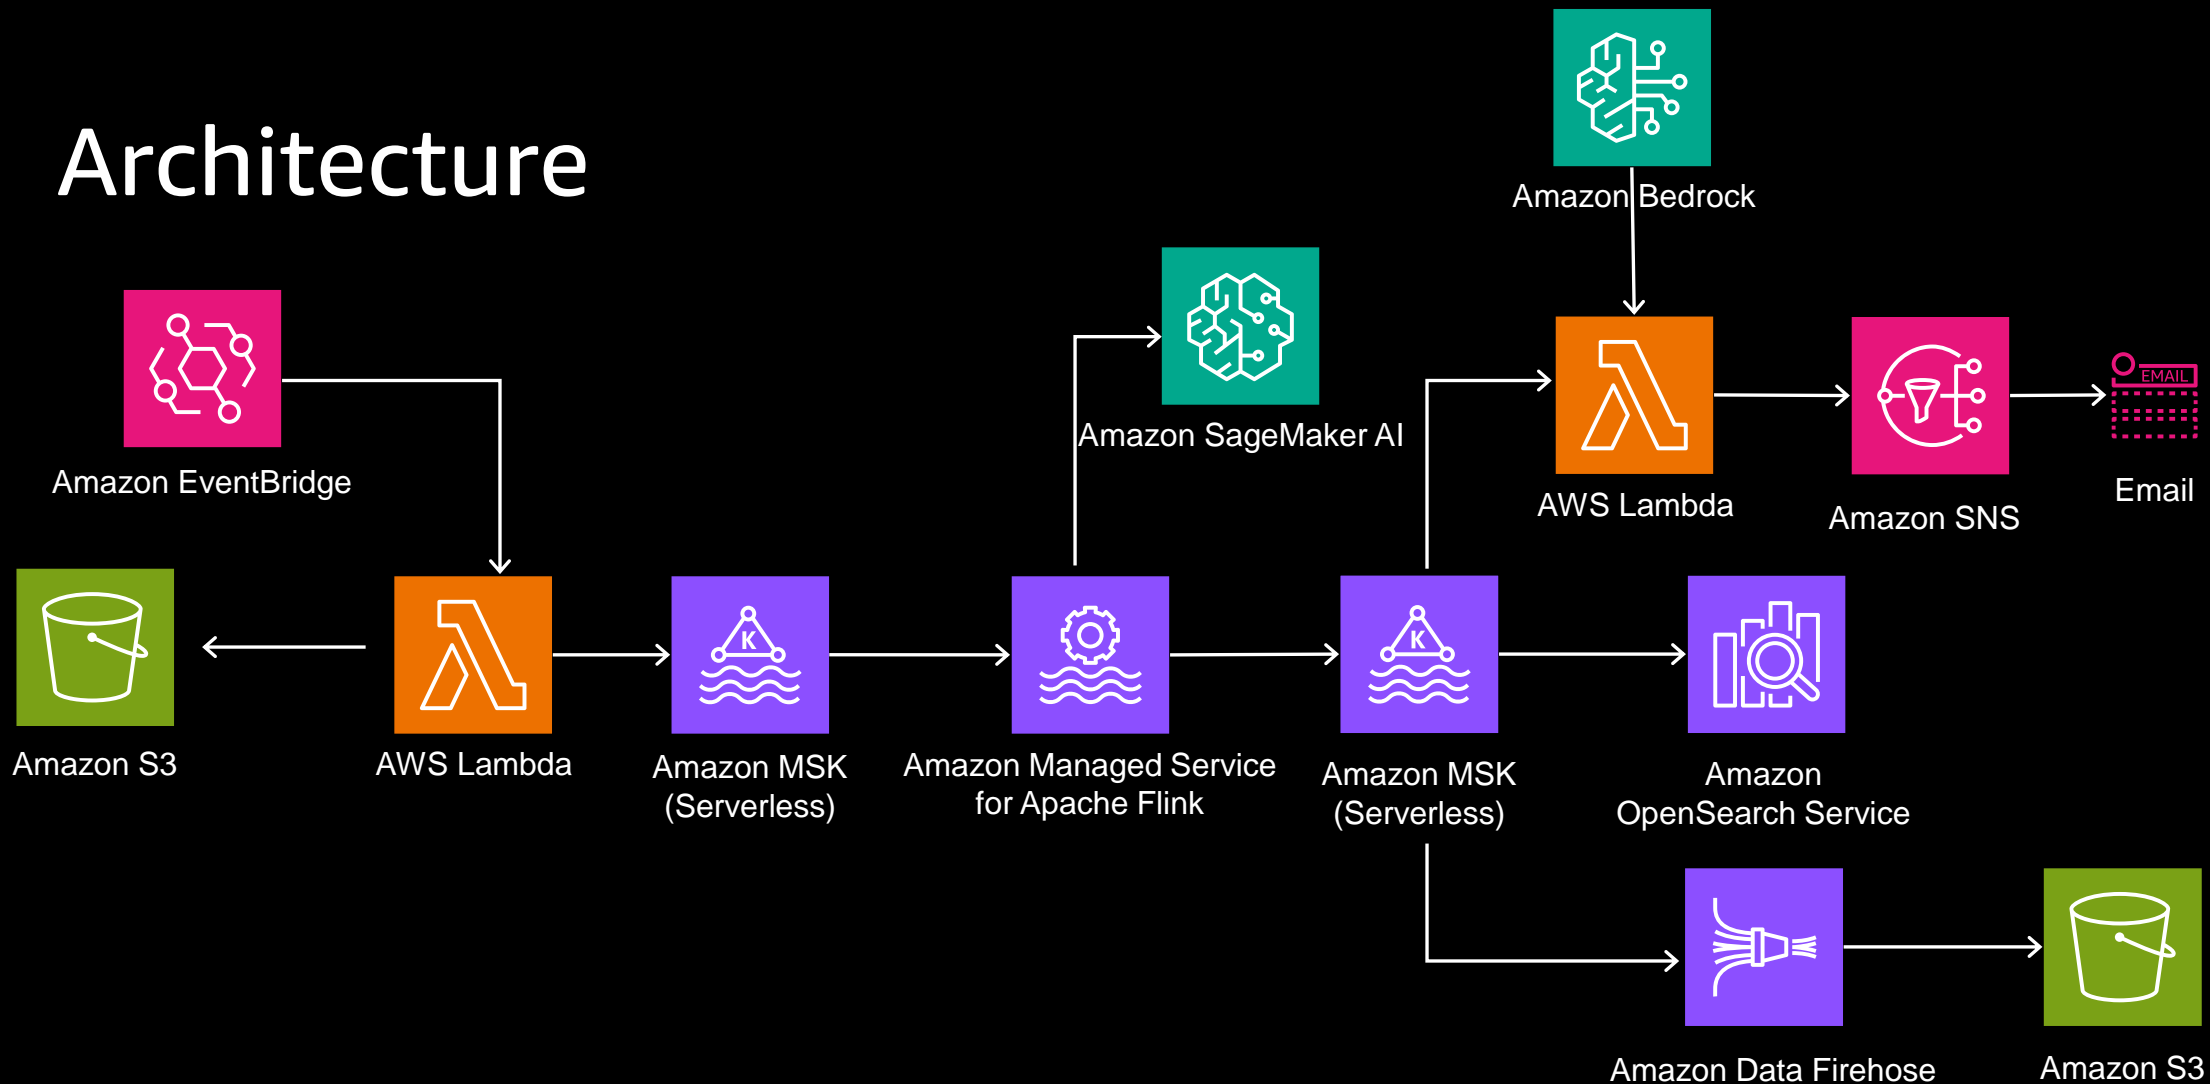
Custom Model Fine-tuning: Customize foundation models with your own data

Scalable Deployment: Deploy generative AI applications at scale

This workshop specifically uses Anthropic *Claude 3 Sonnet v1*



Architecture



Collaboration

Stuck? Don't be afraid to ask your table-mates! Work together to understand and complete the workshop.

Don't hesitate to raise your hand and ask support staff for assistance.



Getting started with Workshop Studio



You have access to an AWS account with everything needed to complete this workshop.



The AWS account is only available for the duration of this session. **You will lose access to the account once the session is complete.**



The workshop environment is deployed to a specific AWS Region. Make sure that you are working in this Region; other Regions are blocked.



Review the terms and conditions of the event. **Do not upload any personal or confidential information to the account.**



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.

Accessing the workshop environment for this session

01

Sign in using your preferred method – <https://join.workshops.aws>

02

Enter the event access code – [\[XXXX-XXXXXX-XX\]](#)

03

Review terms and join event

04

Select **Get Started** to begin the workshop

05

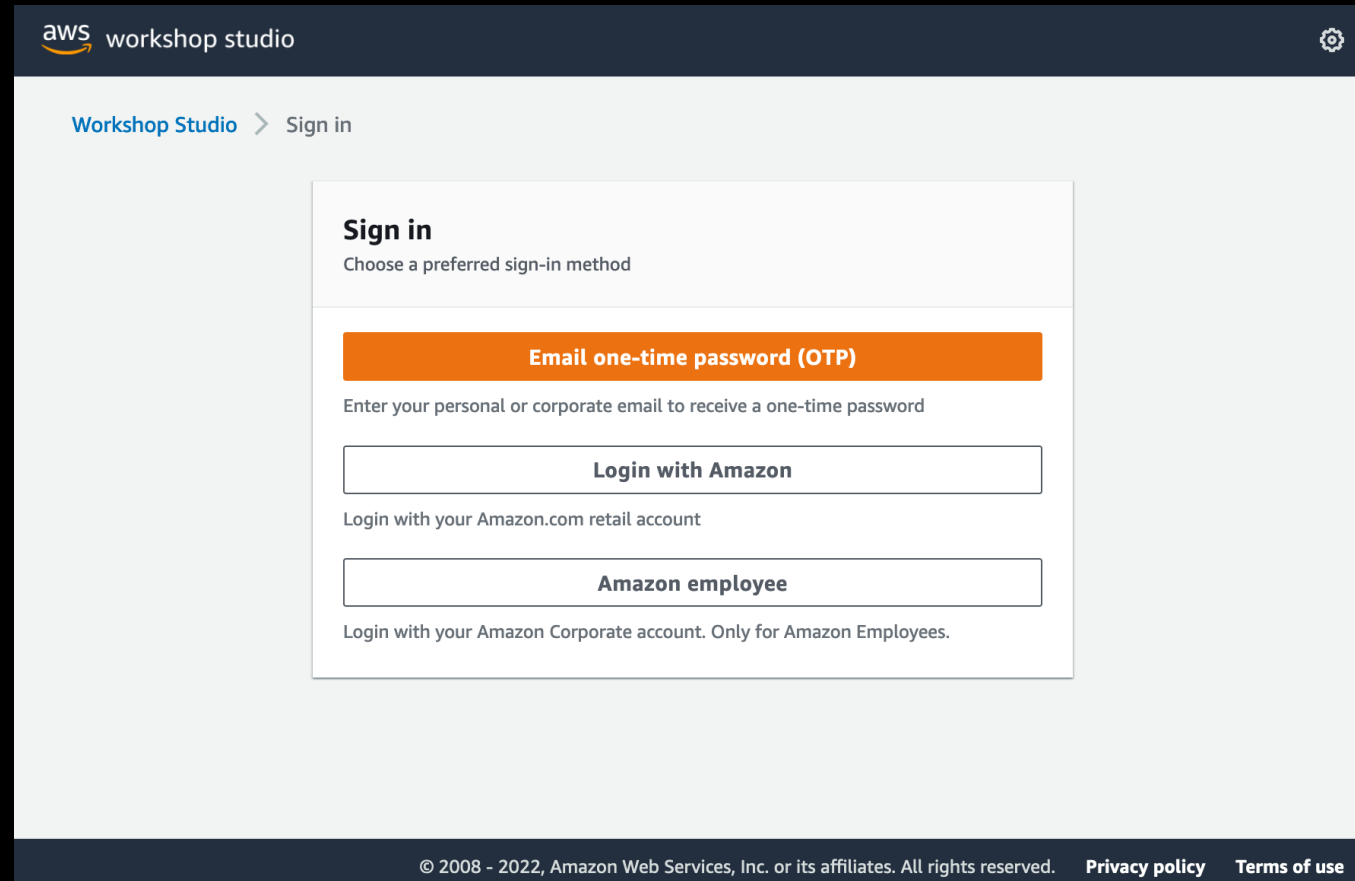
Access AWS accounts – access the AWS Management Console or generate AWS CLI credentials as needed

Next up, hands on Workshop!

- We are going to organize into sections called “checkpoints”
- Before and after each checkpoint, we will review objectives, lessons learned and Q&A
- Feel free to complete at your own pace if you don't want to follow the Checkpoint approach

Hands-On

Workshop Join URL:



The screenshot shows the AWS Workshop Studio sign-in interface. At the top, there's a dark blue header with the 'aws workshop studio' logo on the left and a settings gear icon on the right. Below the header, a breadcrumb trail reads 'Workshop Studio > Sign in'. The main content area is a light gray box with a white border. Inside, the 'Sign in' section is titled 'Sign in' with the subtitle 'Choose a preferred sign-in method'. There are three sign-in options: 1. 'Email one-time password (OTP)' in an orange button, with the instruction 'Enter your personal or corporate email to receive a one-time password'. 2. 'Login with Amazon' in a white button, with the instruction 'Login with your Amazon.com retail account'. 3. 'Amazon employee' in a white button, with the instruction 'Login with your Amazon Corporate account. Only for Amazon Employees.' At the bottom of the page, a dark blue footer contains the copyright notice '© 2008 - 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.' and links for 'Privacy policy' and 'Terms of use'.

aws workshop studio

Workshop Studio > Sign in

Sign in
Choose a preferred sign-in method

Email one-time password (OTP)
Enter your personal or corporate email to receive a one-time password

Login with Amazon
Login with your Amazon.com retail account

Amazon employee
Login with your Amazon Corporate account. Only for Amazon Employees.

© 2008 - 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved. [Privacy policy](#) [Terms of use](#)

Review terms and join event

Workshop Join URL:

aws workshop studio

Workshop Studio > Join event

Step 1
[Enter event access code](#)

Step 2
Review and join

Review and join

Event details

Name	Start time	Duration	Level
AWS General Immersion Day	9/23/2022 01:13 AM	12 hours	-

Description
AWS General Immersion Day

Terms and Conditions

Read and accept before joining the event

1. By using AWS Workshop Studio for the relevant event, you agree to the AWS Event Terms and Conditions and the AWS Acceptable Use Policy. You acknowledge and agree that are using an AWS-owned account that you can only access for the duration of the relevant event. If you find residual resources or materials in the AWS-owned account, you will make us aware and cease use of the account. AWS reserves the right to terminate the account and delete the contents at any time.
2. You will not: (a) process or run any operation on any data other than test data sets or lab-approved materials by AWS, and (b) copy, import, export or otherwise create derivative works of materials provided by AWS, including but not limited to, data sets.
3. AWS is under no obligation to enable the transmission of your materials through Event Engine and may, in its discretion, edit, block, refuse to post, or remove your materials at any time.
4. Your use of AWS Workshop Studio will comply with these terms and all applicable laws, and your access to AWS Workshop Studio will immediately and automatically terminate if you do not comply with any of these terms or conditions.

☒ I agree with the Terms and Conditions

Cancel Previous **Join event**

© 2008 - 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved. [Privacy policy](#) [Terms of use](#)



Access AWS account

Access the AWS Console, or generate AWS CLI credentials as needed



Event ends in 12 hours.

[Event dashboard](#) > Analyzing Network Flow Logs with Generative AI

GenerativeAI-Powered anomaly detection: Spotting anomalies in real-time event

Event information

Start time
9/04/2024 02:31 PM

Duration
12 hours

Accessible regions
us-west-2, us-east-1

Description
Test event for content GenerativeAI-Powered anomaly detection: Spotting anomalies in real-time

Workshop

Title
GenerativeAI-Powered anomaly detection: Spotting anomalies in real-time

Complexity level
400

AWS services
Amazon Bedrock, Amazon Kinesis Data Analytics, Amazon SageMaker

Topics
Analytics, Machine Learning (ML/AI), Generative AI

Description
Uncover the intricacies of analyzing indefinite log data events with AI-Powered threat and anomaly detection. Explore the dynamic synergy of data streaming and Splunk for real-time threat and anomaly detection. Coupled with the integration of generative AI, this talk provides adaptive defense strategies against evolving security threats.

[Get started >](#)

GenerativeAI-Powered anomaly detection: Spotting anomalies in real-time event


[Analyzing Network Flow Logs with Generative AI](#)
[Use Case Overview](#)
[Data Streaming with AI](#)
[Launch Resources](#)
[Data Ingestion with Lambda and Amazon MSK](#)
[Examine Data in MSF Studio](#)
[Process Data in Real-Time](#)
[View Anomalous Output using MSF Studio](#)
[Create Smart Alerts with Gen AI](#)
[Configure Amazon Data Firehose](#)
[Examine Results and Cleanup](#)

AWS account access
[Open AWS console \(us-west-2\)](#)
[Get AWS CLI credentials](#)
[Get EC2 SSH key](#)

Event ends in 11 hours 58 minutes.

[Event dashboard](#) > Analyzing Network Flow Logs with Generative AI

Analyzing Network Flow Logs with Generative AI



Checkpoint 1

Analyzing Network Flow Logs with Generative AI


Launching Resources

- Use Case Overview
- Data Streaming with AI
- Data Ingestion with Lambda and Amazon MSK
- Examine Data in MSF Studio
- Process Data in Real-Time
- View Anomalous Output using MSF Studio
- Create Smart Alerts with Gen AI
- Configure your MSK Serverless Consumers
- ✳ Where to go from here?
- Examine Results and Cleanup


Analyzing Network Flow Logs with Generative AI

Analyzing Network Flow Logs with Generative AI


Checkpoint 1



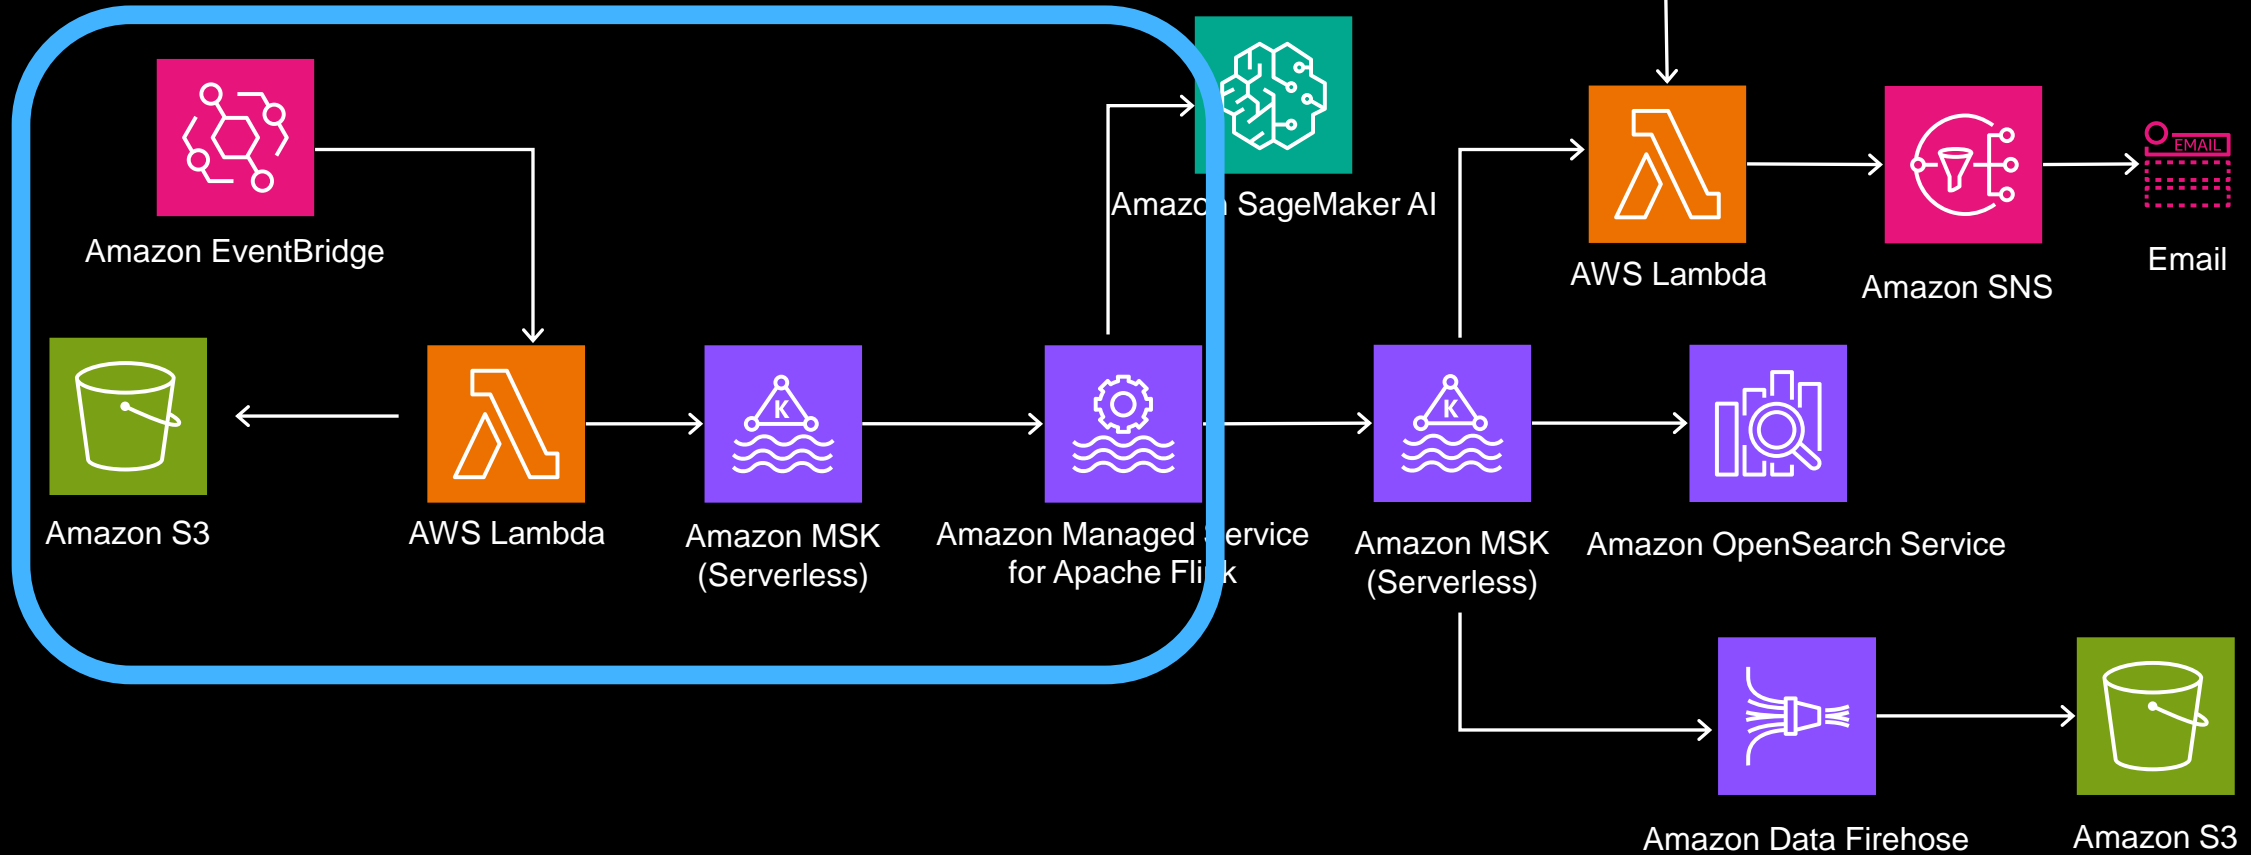
Checkpoint 2



Checkpoint 3



Checkpoint 1



Finish section “Examine in MSF Studio”

Checkpoint 1 Review

- Producing Simulated Data to MSK Cluster via EventBridge Rule with Lambda
- We can validate / see that data in MSF Studio in real-time

Checkpoint 2

Analyzing Network Flow Logs with Generative AI

Launching Resources

Use Case Overview

Data Streaming with AI

Data Ingestion with Lambda and Amazon MSK

Examine Data in MSF Studio

Process Data in Real-Time

View Anomalous Output using MSF Studio

Create Smart Alerts with Gen AI

► Configure your MSK Serverless Consumers

✳ Where to go from here?


Examine Results and Cleanup

Analyzing Network Flow Logs with Generative AI

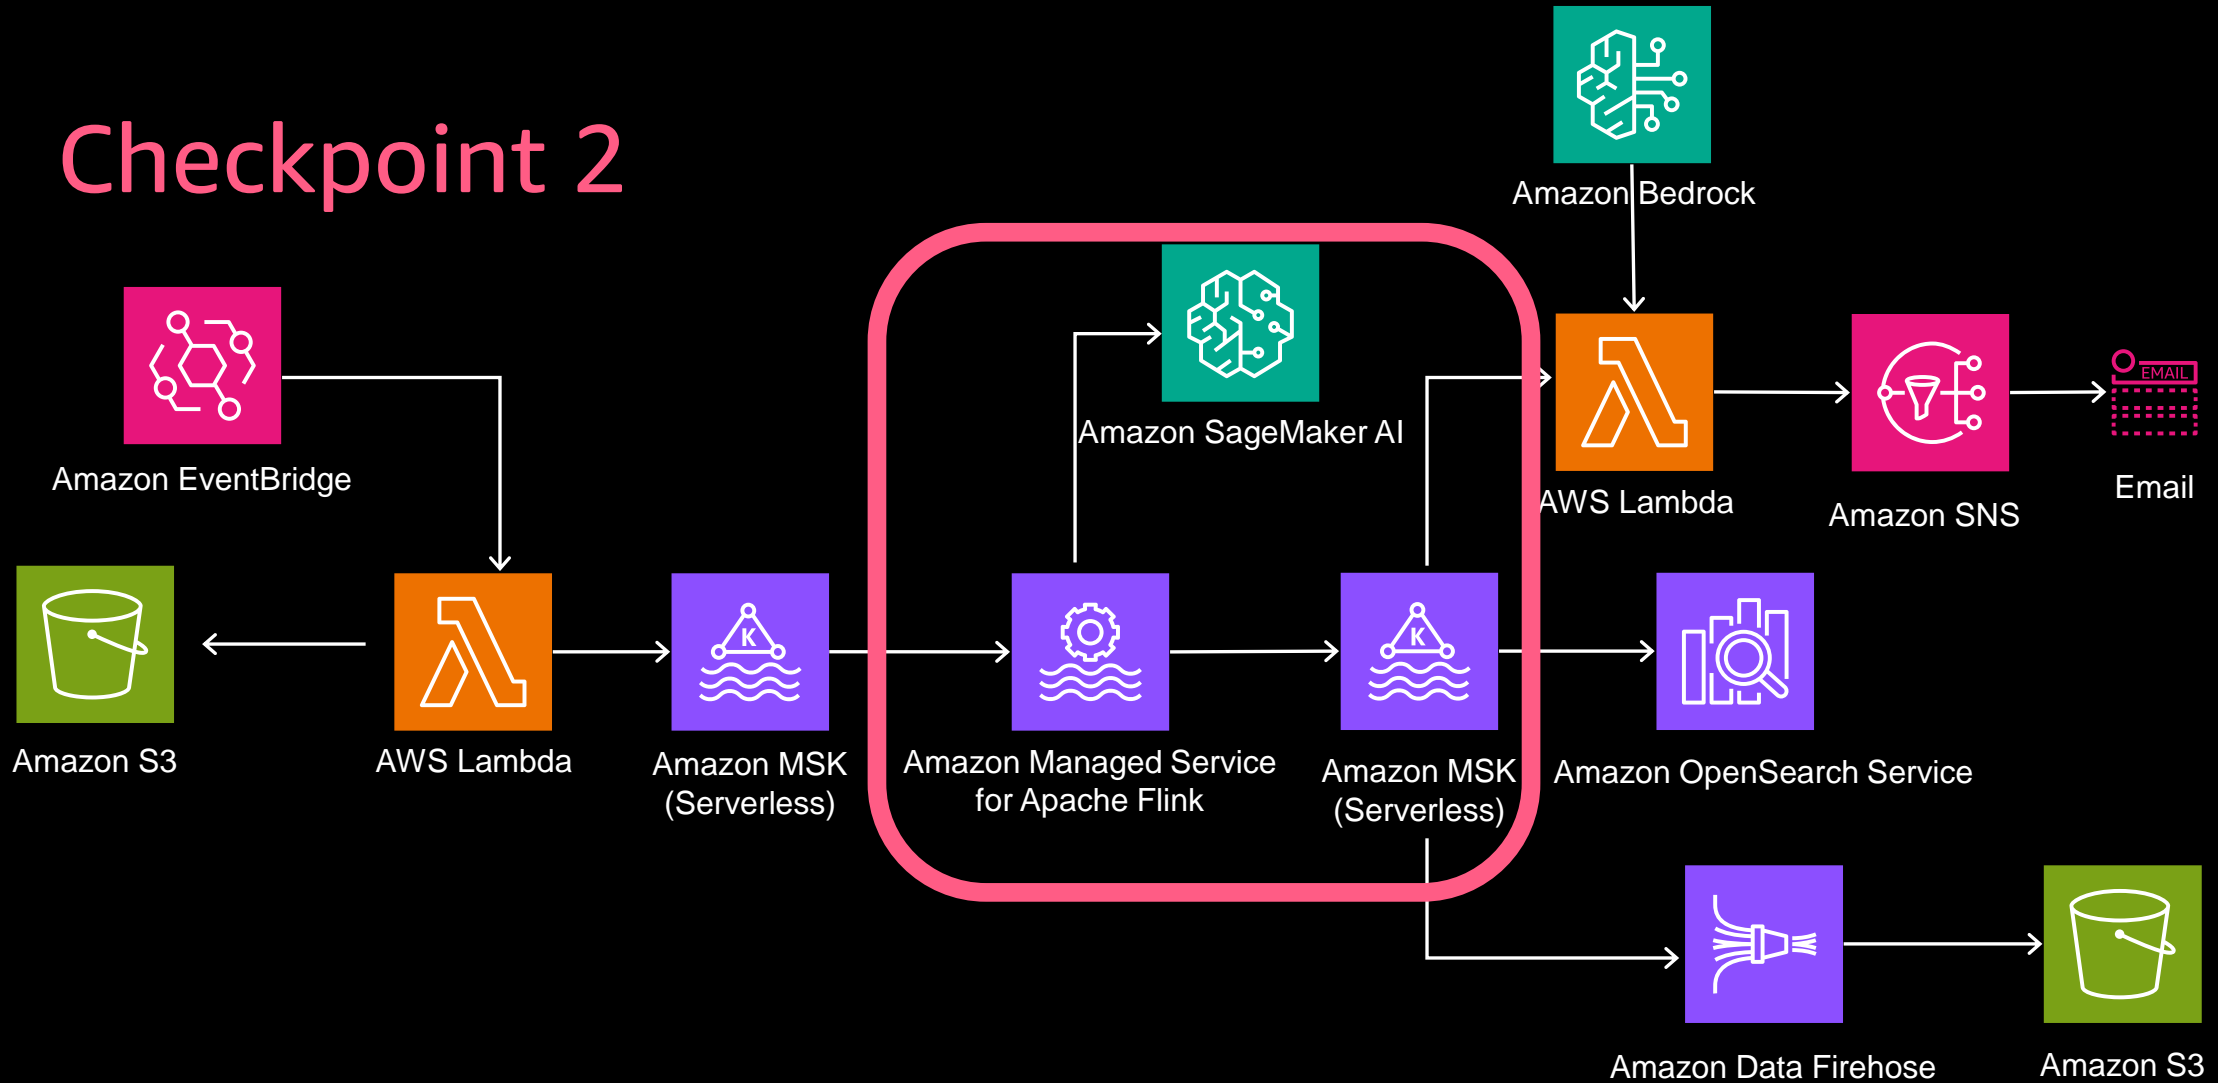
Analyzing Network Flow Logs with Generative AI



Checkpoint 2



Checkpoint 2



Finish Section “View Anomalous Output using MSF Studio”

Checkpoint 2 Review

- Validated our model in Sagemaker AI
- Started a durable MSF Application with state that queries SageMaker AI in real-time
- Validated we can view the output of MSF application in MSF Studio
- All the while interacting with MSK Serverless

Checkpoint 3

Analyzing Network Flow Logs with Generative AI

Launching Resources

Use Case Overview

Data Streaming with AI

Data Ingestion with Lambda and Amazon MSK

Examine Data in MSF Studio

Process Data in Real-Time

View Anomalous Output using MSF Studio

Create Smart Alerts with Gen AI

► Configure your MSK Serverless Consumers

✿ Where to go from here?


Examine Results and Cleanup

Analyzing Network Flow Logs with Generative AI

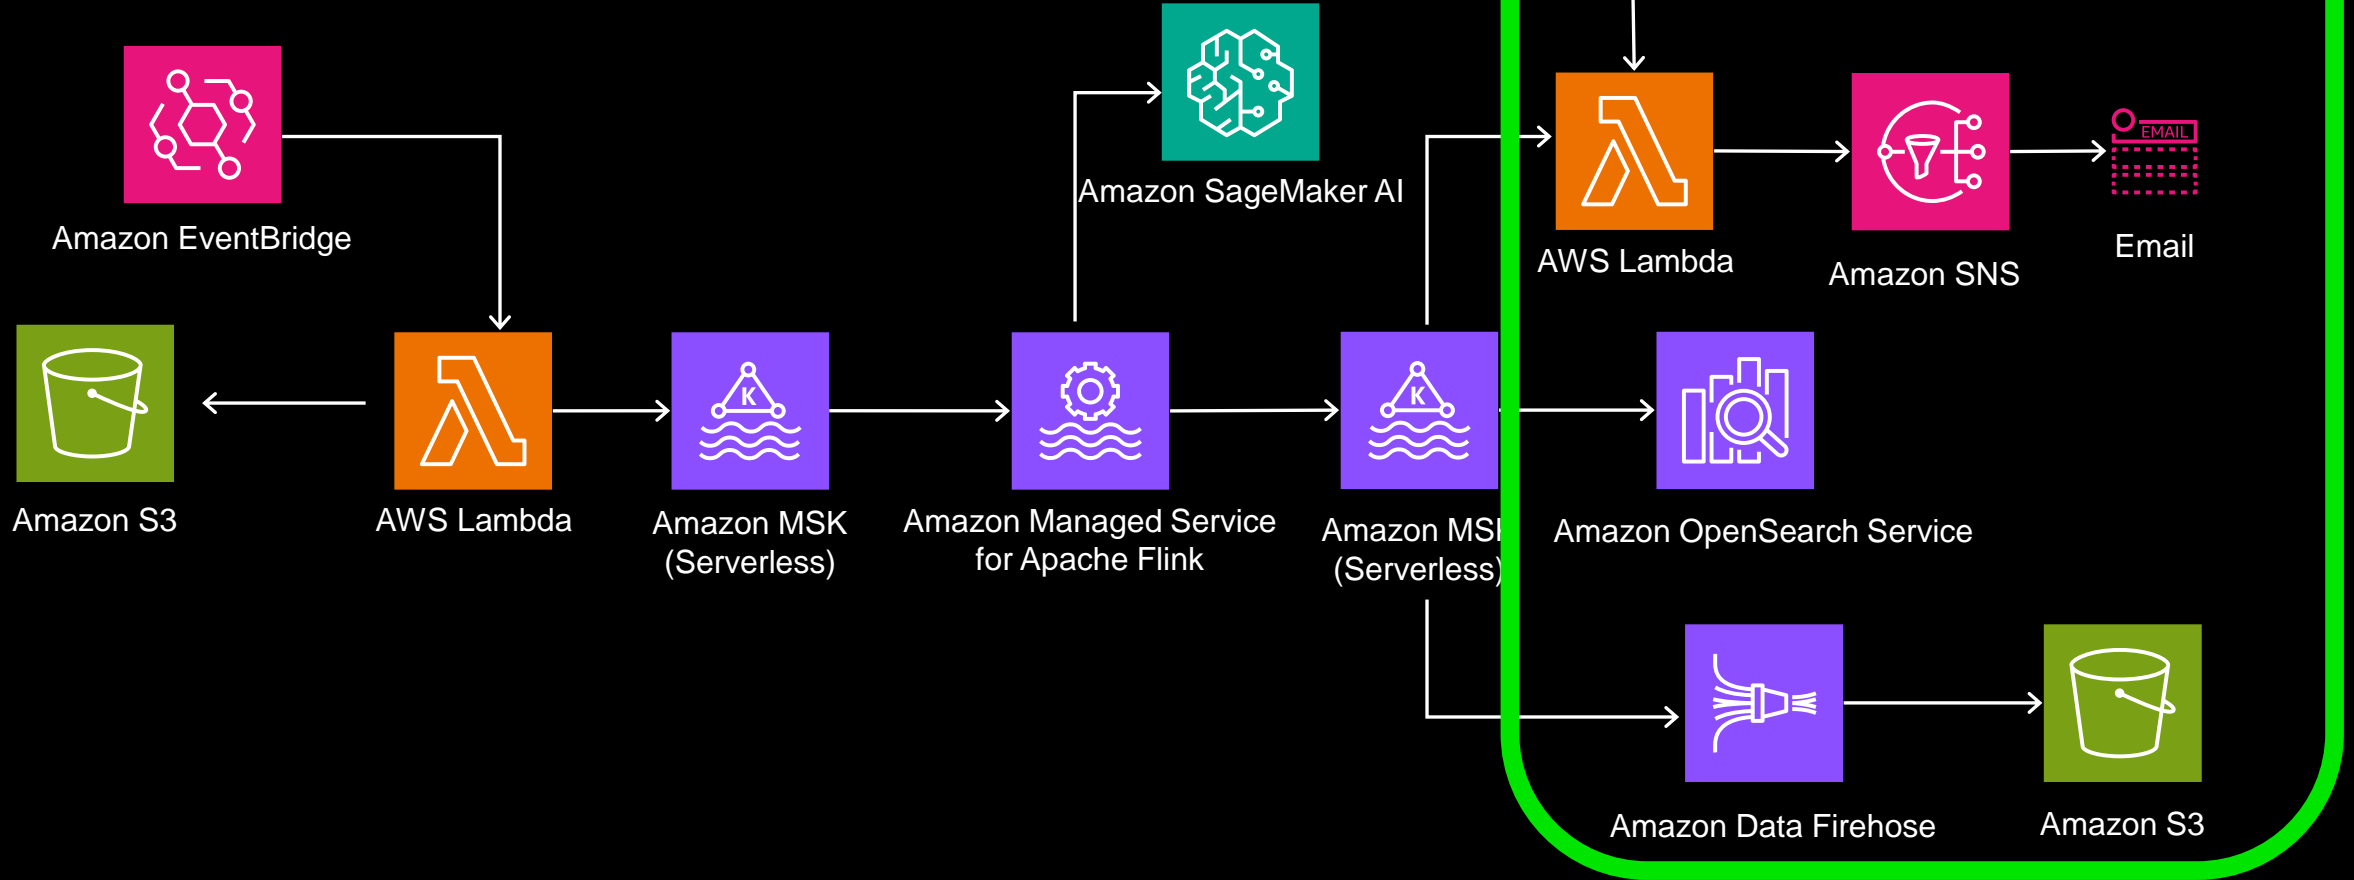
Analyzing Network Flow Logs with Generative AI



Checkpoint 3



Checkpoint 3



Finish Section “Configure your MSK Serverless Consumers”

Checkpoint 3 Review

- Explored Amazon Bedrock with anomalous data
- Triggered and received e-mails related to anomalies
- Ingested data into Amazon Opensearch and explored in Dashboard
- Backed data up using Amazon Data Firehose

Extra Credit!

If you've completed Checkpoints 1, 2 and 3, feel free to continue on with the workshop, working through some challenges we've added to dive even deeper

Thank you!

Joe Khazen

Principal, WW GTM
Specialist, Data
Streaming (AWS)

Austin Groeneveld

Assoc. WW SSA Data
Streaming (AWS)

Ali Alemi

Sr. WW SSA Data
Streaming (AWS)

